



## **Truth Behind DeepFakes(TBD) GROUP - 9**

# **CS 776 : Deep Learning For Computer Vision**

Department of Computer Science and Engineering  
Indian Institute of Technology Kanpur (IITK)

# Expected Outcome and deliverables from Initial Presentation

- Taking computational capacity into account, we will design a ~~Streamlit~~ Gradio based frontend for the detection of Deep Fake images and videos.
- **Phase 1: Deep Fake Image Detection**
  - Utilize a CNN and train a classifier on the publicly available 140k Real and Fake Faces dataset from Kaggle. (~~DenseNET/VGGFace/XceptionNet/Custom GNN Architectures~~/Efficientnet\_b0).
  - Conduct robustness analysis to further evaluate the model.
- **Phase 2: Deep Fake Video Detection**
  - Initially employ the ~~mini\_face\_forensics~~ used original FF++ compressed c23 LQ dataset from Kaggle for model training and evaluation.
  - We'll explore multiple model architectures to capture both spatial and temporal features.:
    - CNN + LSTM combination.
    - ViTs, such as ~~GViT~~/Swin Transformer/~~FasterViT~~
    - Model Ensembling
  - Extract features using state-of-the-art pretrained models such as ~~InceptionV3 or XceptionNet~~ used Efficientnet\_b0 and ResNext50 for transfer learning. The pretrained model will be used to obtain a feature vector, further the LSTM layers will be trained using these features and create a baseline model.
- Implementation Strategy:
  - Integrate various techniques from research papers to improve baseline model performance.
  - If feasible, apply the finalized model to a state-of-the-art Faceforensics++ dataset for further validation.
- Deliverables:
  - A web-based interface for Deep Fake detection with Python scripts.
  - Report documenting the implementation.

# Deepfake Image Detection

- Deepfake technology is advancing rapidly, posing serious threats to media trust and security.
- Build a model to classify fake image as Real or Fake

## Dataset Details

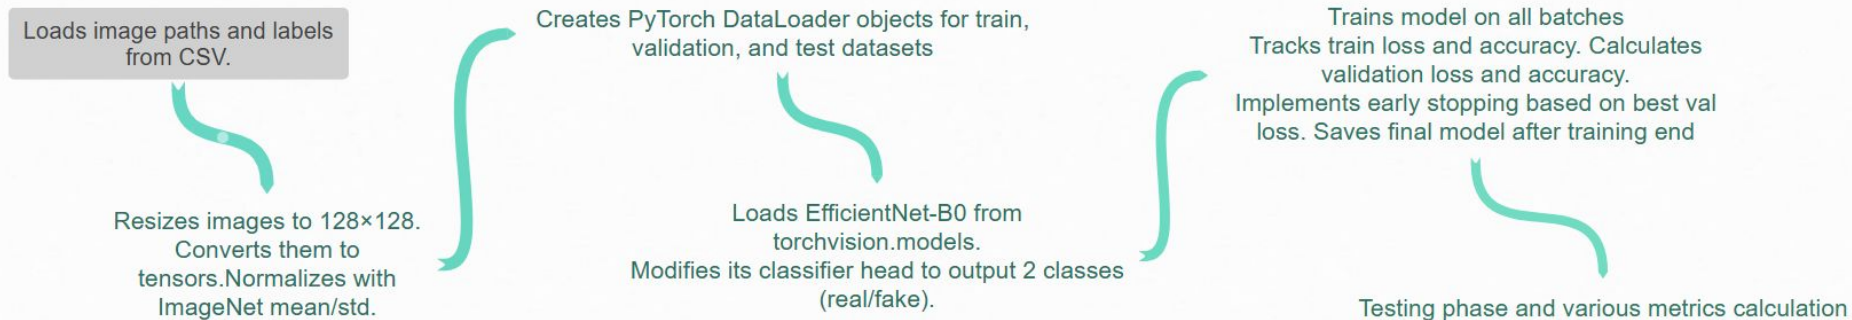
- **Source:** 140K real and fake face images from Kaggle
- **Subset Used:** 80K images total
  - Training : We have used 30k Fake and 30k Real Image.
  - Validation : 10K images
  - Testing : 10k images

## Preprocessing Pipeline


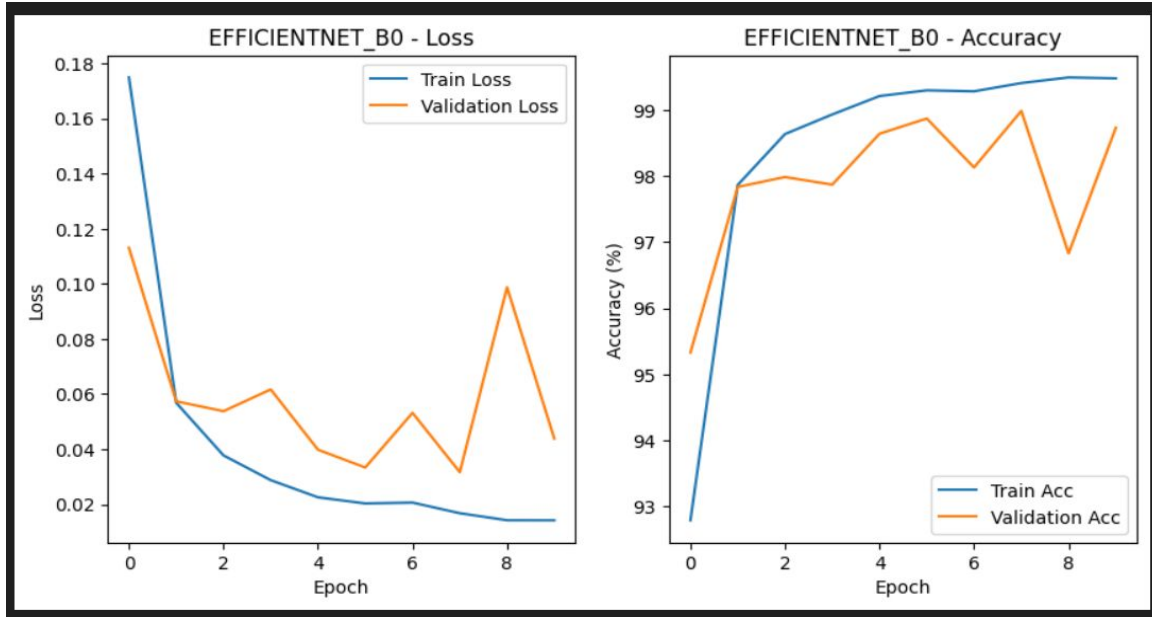
- To ensure consistency and optimize model performance, the following preprocessing steps were applied
- **Resizing** : All images were resized to a fixed dimension of **128 × 128 pixels** to maintain uniformity in input size.
- **Normalization** : Pixel values were normalized using ImageNet statistics
  - Mean [0.485, 0.486, 0.406]
  - Standard Deviation : [0.229, 0.224, 0.225]
  - This helps align the data distribution with what the pretrained EfficientNet-B0 model expects.

# Training Configuration

- Base Model: EfficientNet-B0  
Pretrained on ImageNet (1M+ images, 1000 classes)
- Epochs: 10
- Batch Size: 64
- Loss Function: CrossEntropyLoss
- Optimizer: Adam (lr = 0.001)
- Precision: Mixed-Precision Training (torch.amp)
- Device: GPU (CUDA-enabled if available)



# Results & Performance



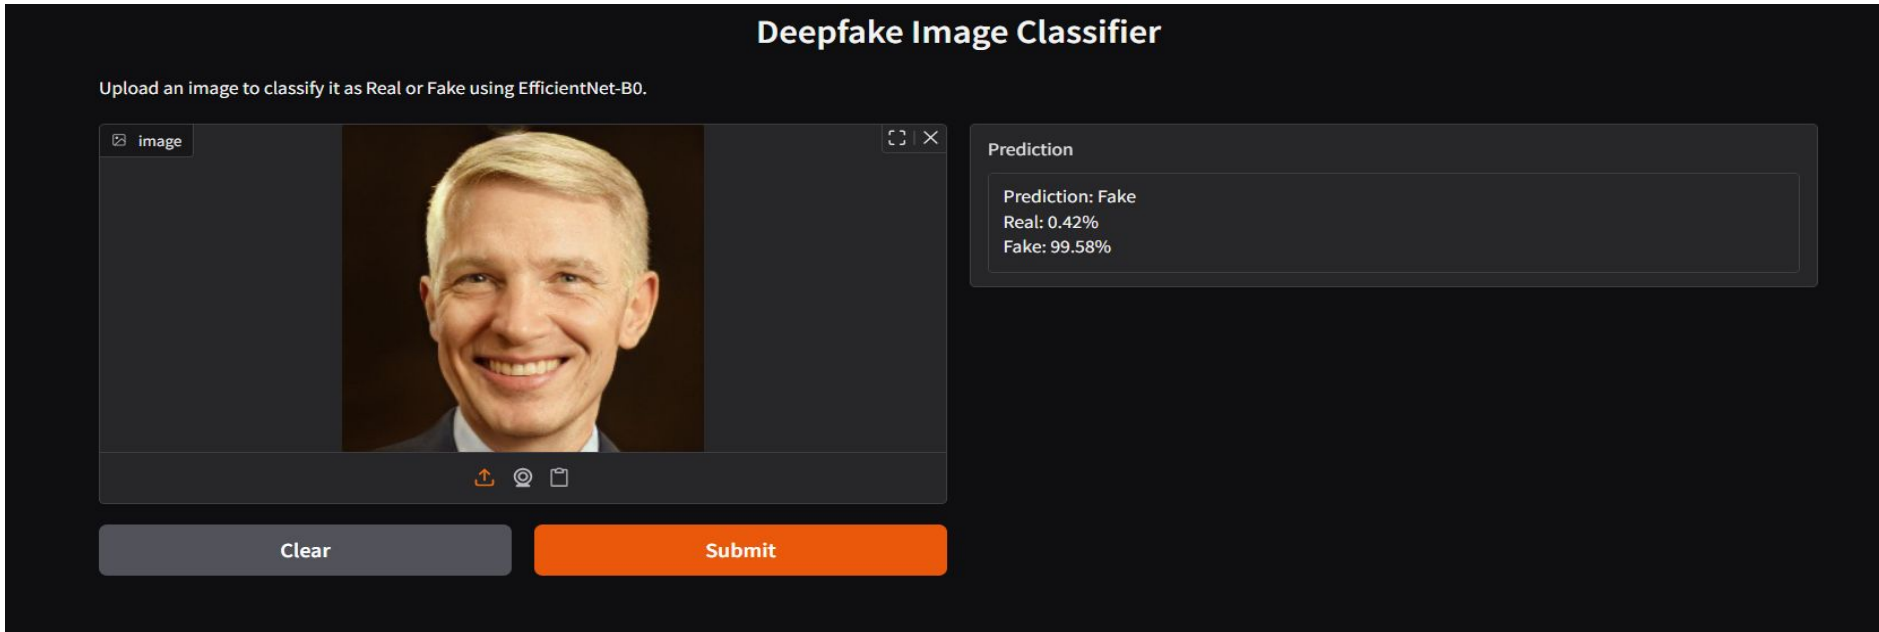
Classification Report:

	precision	recall	f1-score	support
Fake	0.99	0.99	0.99	10000
Real	0.99	0.99	0.99	10000
accuracy			0.99	20000
macro avg	0.99	0.99	0.99	20000
weighted avg	0.99	0.99	0.99	20000

Plotted **loss and accuracy curves** across 10 epochs to monitor learning behavior

# Interface for Real/Fake Prediction

- Gradio provides a simple and intuitive web interface for testing our deepfake classification model in real-time.
- Upload image → Model predicts: Real / Fake



# Video Deepfake Detection

Pretrained CNN's +LSTM:

- ResNeXt50/EfficientNet-B0 as feature extractor.
- LSTM for temporal aggregation.

Swin Transformer:

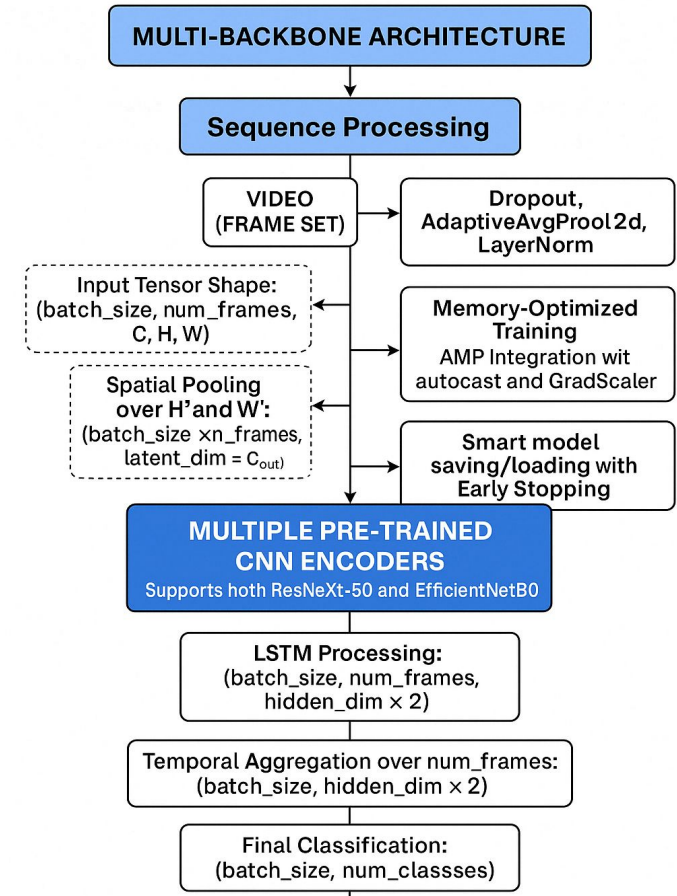
- Patch-based shifted attention for spatial features.
- Sequence averaging for temporal consistency.

Dataset Used for Train/Test/Val: FaceForensics++ c23 (LQ)

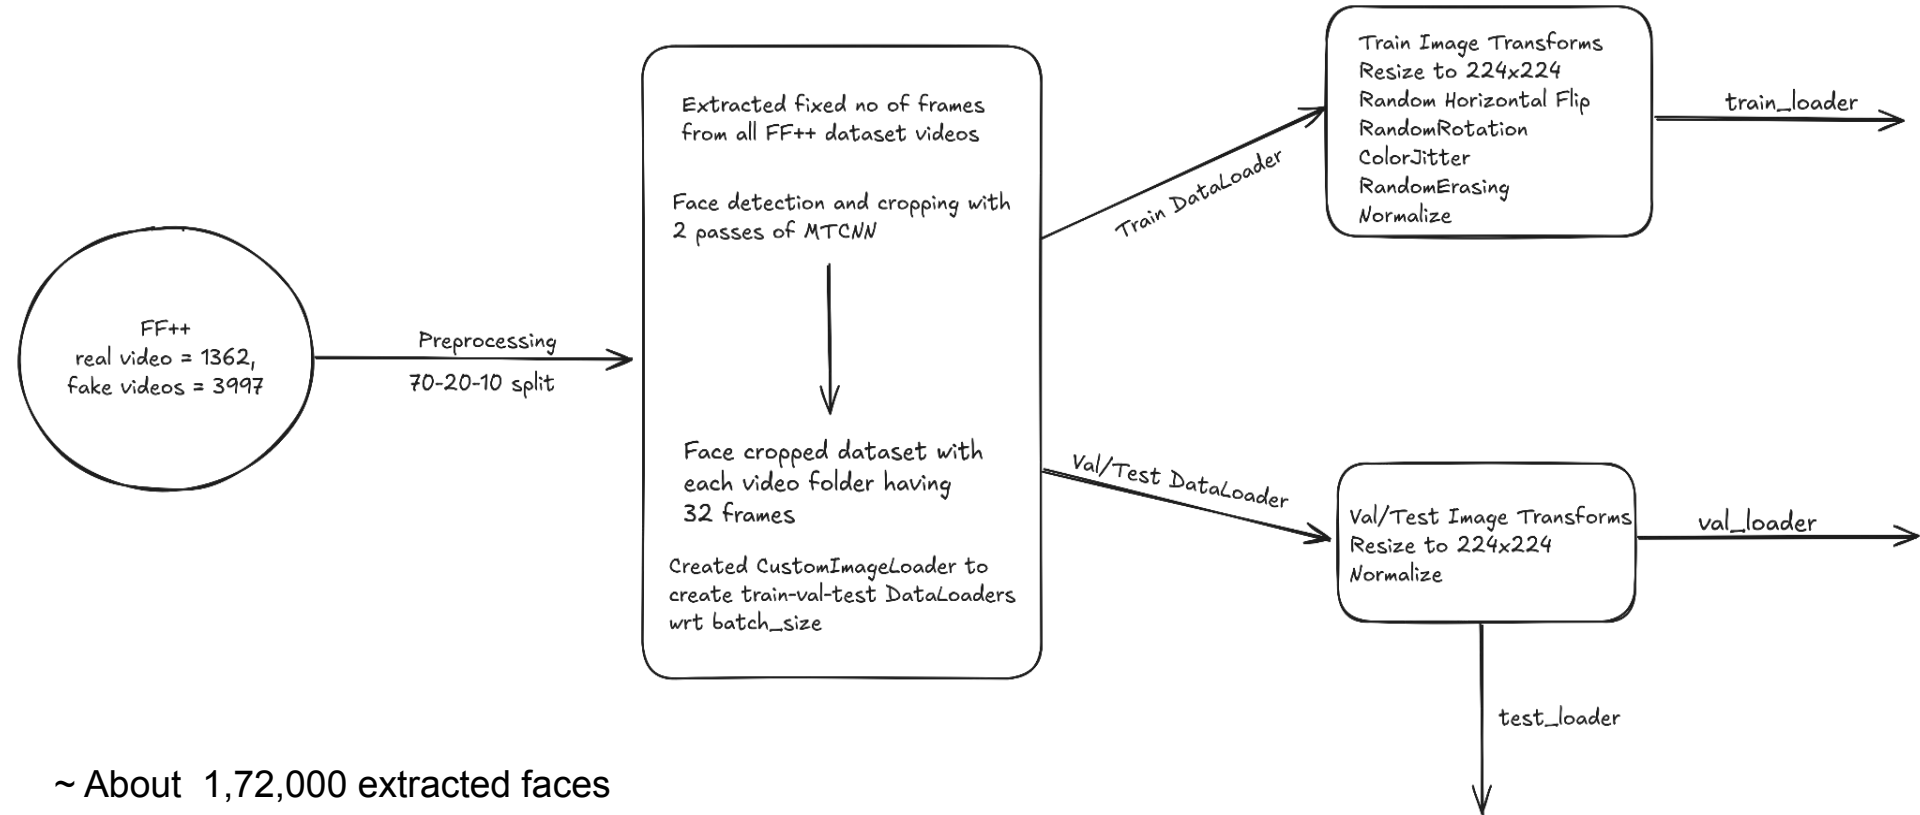
- Real
  - YT 1000 videos
  - Actor 363 videos
- Fake
  - Deepfakes 1000
  - Face2Face 1000
  - FaceShift 1000

Generalization Dataset:

- CelebDFv1
- UADFV Dataset



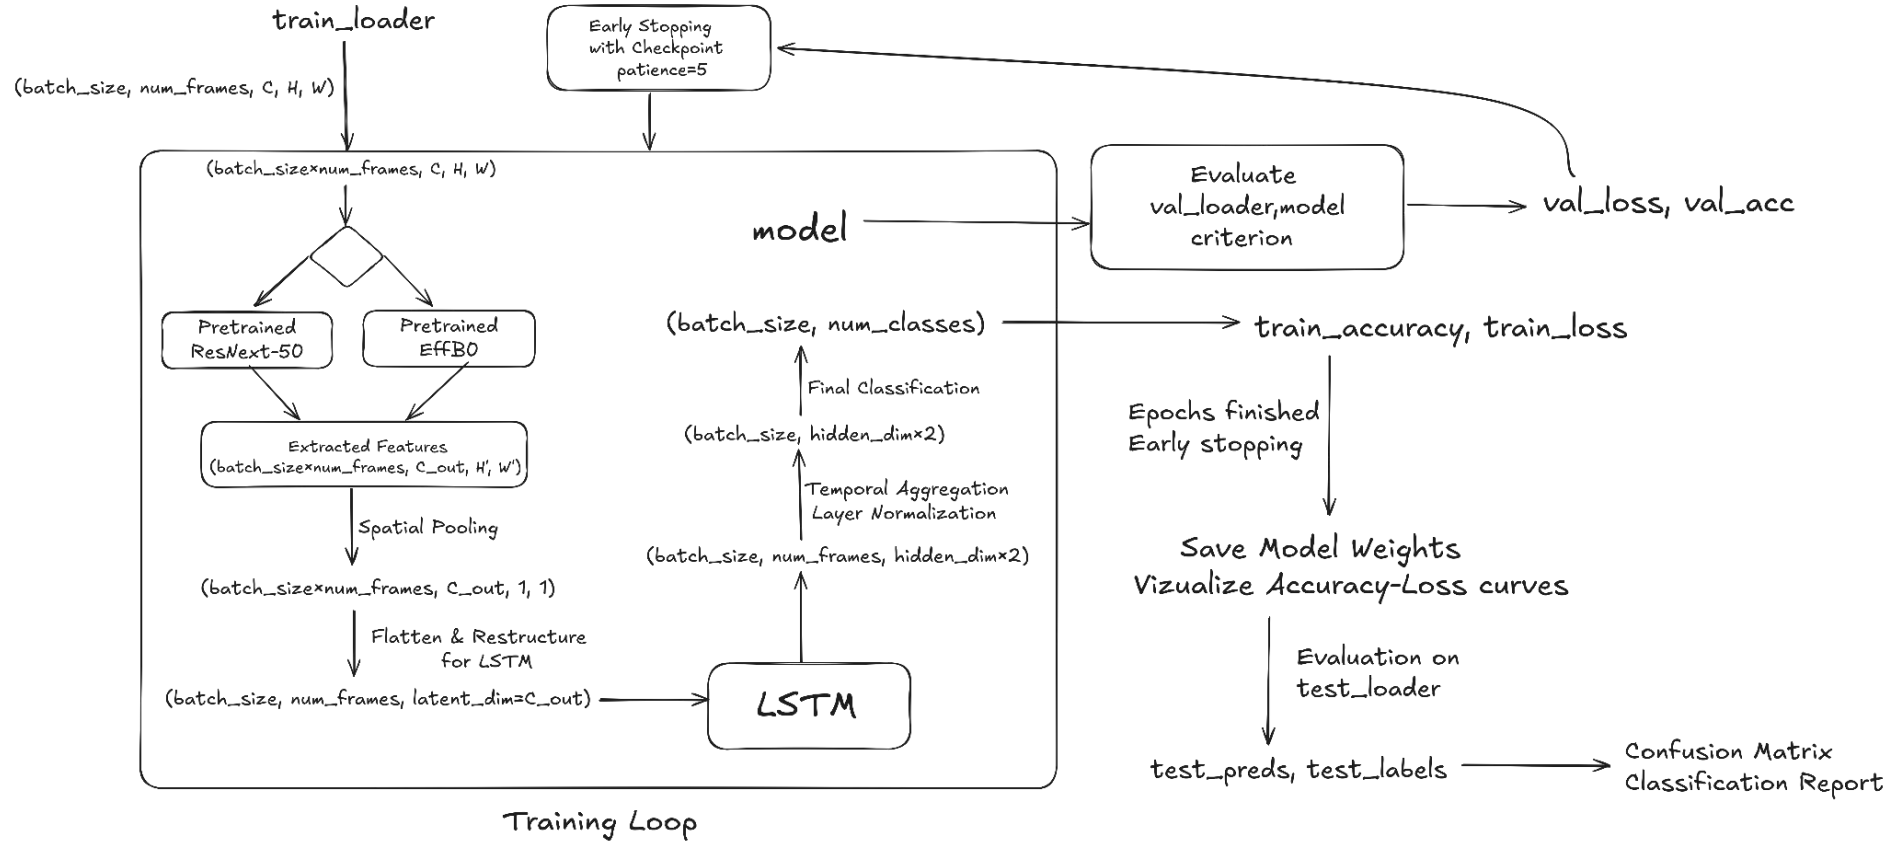
# Preprocessing



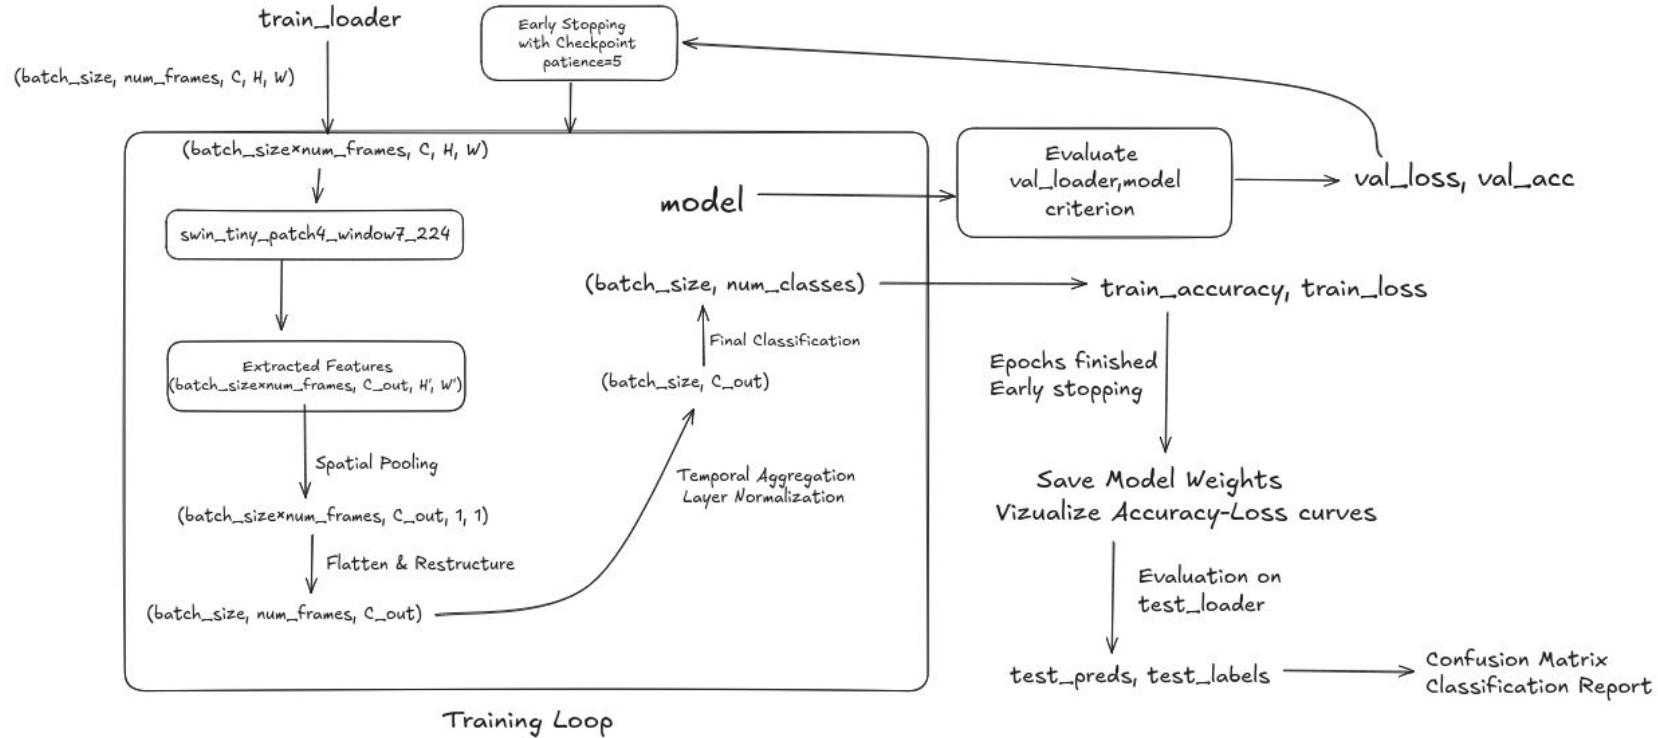
~ About 1,72,000 extracted faces



# Training Workflow Pipeline I



# Training Workflow Pipeline II



# Project novelty

- Dual-Stage Face Verification
  - MTCNN used twice: First for detection, second for verification
  - **Why Better:** Reduces false face detections that plague many deepfake detectors. Cleaner input data than papers using single-pass detection
- Frame Sampling Strategy
  - Linear interpolation-based uniform sampling across entire video duration
  - **Why Better:** Avoids bias toward initial frames seen in many papers that simply take first N frames
  - **Benefit:** Captures temporal patterns from the entire video timeline. Dynamically adapts to videos of any length.
- Extreme Augmentation Pipeline
  - Aggressive train-time augmentations for robustness not commonly used in video deep fake detection
  - **Why Better:** Simulates real-world distortions better than basic flips/rotations.
- Optimized Class Balancing
  - Calculated class weights Based on dataset statistics for CE loss  $(N_{\text{total}}/2)/N_{\text{real}}$
  - **Why Better:** More mathematically grounded than papers using naive sampling.
- Multi-Backbone Architecture and Hybrid Normalization Strategy
  - Combined Spatial Averaging (CNN) and temporal normalization(LSTM)
  - **Why Better:** Handles variance in both facial regions and temporal dynamics. Most papers only use BatchNorm. Combines spatial (channel-wise) and temporal (sequence-wise) normalization.

# Project novelty

## • Memory-Optimized Sequence Processing with LSTM

- AMP (Automatic Mixed Precision) integration: autocast and GradScaler enabling 16-bit training without gradient underflow in temporal layers. Seeding numpy random, PyTorch, torch.cuda.manual\_seed\_all etc.
- **Why Better:** Enables larger batch sizes and systematically eliminates common reproducibility pitfalls by seeding
- **Benefit:** 2-3x faster training than typical implementations



- This architecture efficiently bridges 2D CNNs (spatial processing) and LSTMs (temporal modeling)..
- Total Time Complexity (per forward pass) CNN-LSTM: Total Floating Point Operations (FLOPs) per video

$$\mathcal{T}_{\text{total}} = \underbrace{B \cdot F \cdot \mathcal{T}_{\text{CNN}}}_{\text{CNN Cost}} + \underbrace{B \cdot F \cdot \mathcal{T}_{\text{LSTM}}}_{\text{LSTM Cost}}$$

$$\mathcal{T}_{\text{CNN}} = \sum_{\ell=1}^D \underbrace{C_{\text{in}}^{\ell} \cdot C_{\text{out}}^{\ell} \cdot k_{\ell}^2 \cdot H_{\ell} \cdot W_{\ell}}_{\text{Layer } \ell}$$

$$\mathcal{T}_{\text{LSTM}} = 4 \cdot \underbrace{h^2}_{\text{Hidden-Hidden}} + \underbrace{h \cdot C_{\text{out}}}_{\text{Input-Hidden}}$$

$$\mathcal{T}_{\text{total}} = B \cdot F \cdot \left( \sum_{\ell=1}^D C_{\text{in}}^{\ell} C_{\text{out}}^{\ell} k_{\ell}^2 H_{\ell} W_{\ell} + 4h(h + C_{\text{out}}) \right)$$

B: Batch Size

F: No of frames used per video

D: CNN depth

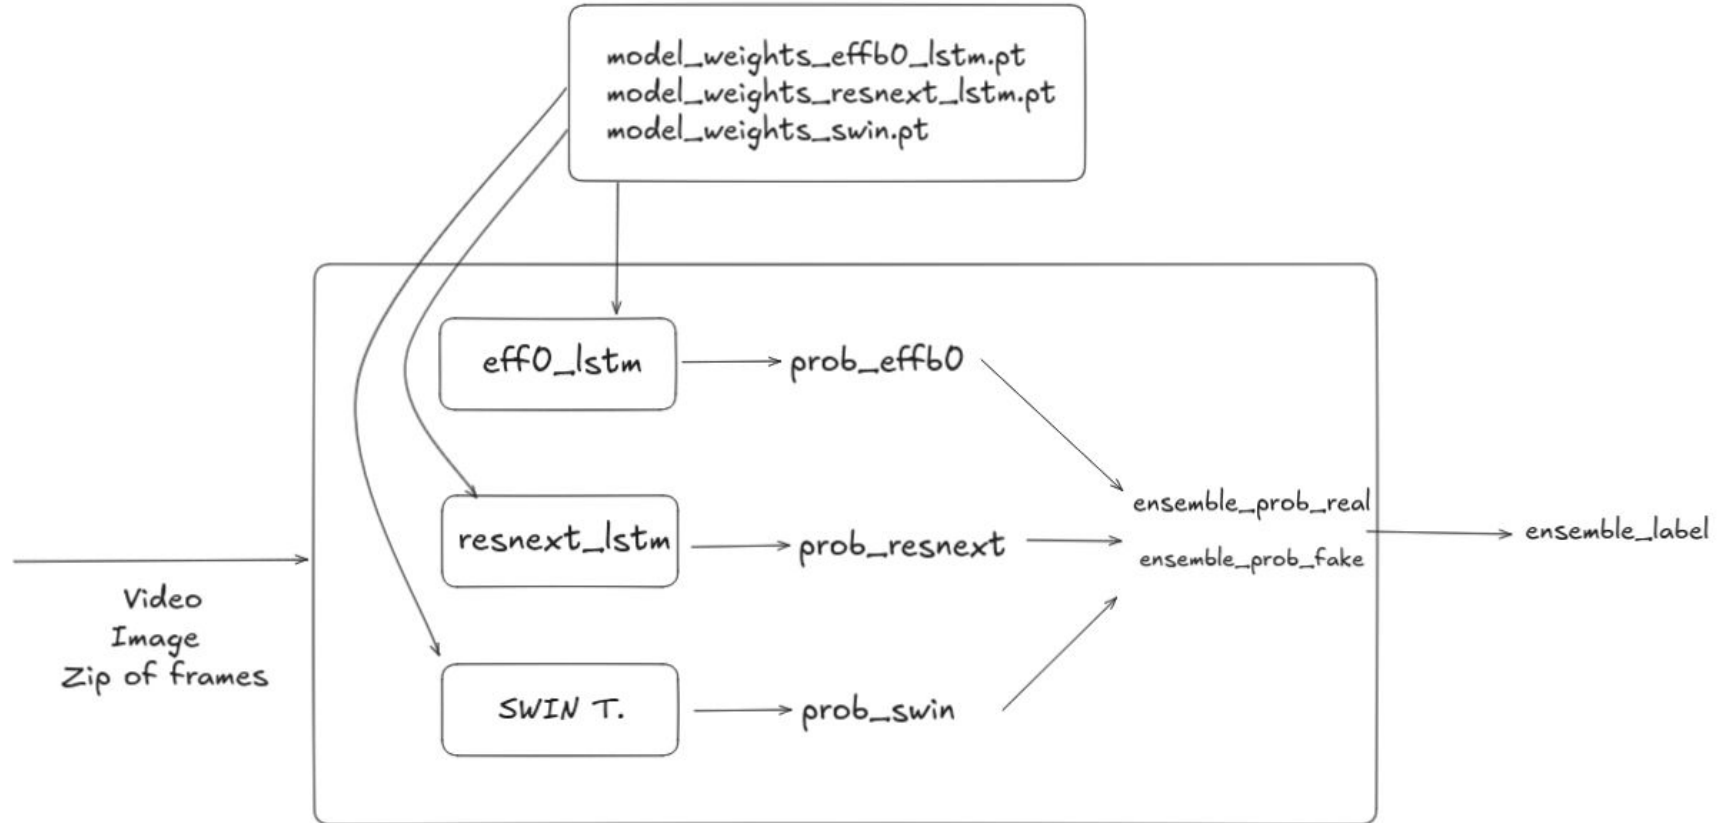
h: LSTM hidden units

C<sub>in</sub><sub>ℓ</sub>: Input Channels for Layer ℓ

C<sub>out</sub><sub>ℓ</sub>: Output Channels for Layer ℓ

C<sub>out</sub>: CNN Output Features

# Model Ensembling/Prediction Workflow



# Training Configuration

The experiments were conducted on an NVIDIA RTX 4060 Laptop GPU, which has 8GB of VRAM.

For SWIN Transformer, experiments were done on Kaggle T4 GPU with 15GB of VRAM.

Parameter	Value
Batch Size	2/4/8
Number of frames	10/20/30
Epochs	20
Optimizer	AdamW
Learning Rate	1e-4
Scheduler	Lr_scheduler step_size 5, gamma 0.5
Loss Function	Weighted CrossEntropyLoss
Activation Function	LeakyRelu
Early Stopping Patience	5 epochs
Dropout Rate	0.4




Extensively used CUDA with consistent GPU utilization (98-100%).

Critical VRAM overflow at sequence length 30 on RTX4060 with batch size = 4.

Model	#Parameters	Size
Effb0+Lstm	5.583M	23MB
ResNext+Lstm	25.342M	102MB
Swin	27.52M	110MB

# Experimental Results- Efb0+Lstm

BS: 4  
F: 30




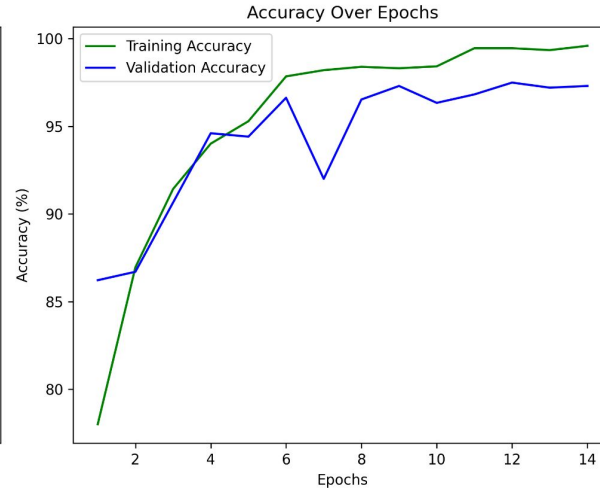

	precision	recall	f1-score	support
Real	0.9440	0.9219	0.9328	128
Fake	0.9749	0.9823	0.9786	396
accuracy			0.9676	524
macro avg	0.9595	0.9521	0.9557	524
weighted avg	0.9674	0.9676	0.9674	524

Training finished! Total Time: 153.738 min

BS=8	F=10	94.79%
BS=4	F=20	96.08%
BS=4	F=30	96.76%

BS: 2  
F: 30

# Experimental Results- Resnext50+Lstm



	precision	recall	f1-score	support
Real	0.9104	0.9531	0.9313	128
Fake	0.9846	0.9697	0.9771	396
accuracy			0.9656	524
macro avg	0.9475	0.9614	0.9542	524
weighted avg	0.9665	0.9656	0.9659	524

BS = 8	F=10	95.16%
BS = 4	F=20	95.52%
BS = 2	F=30	96.56%

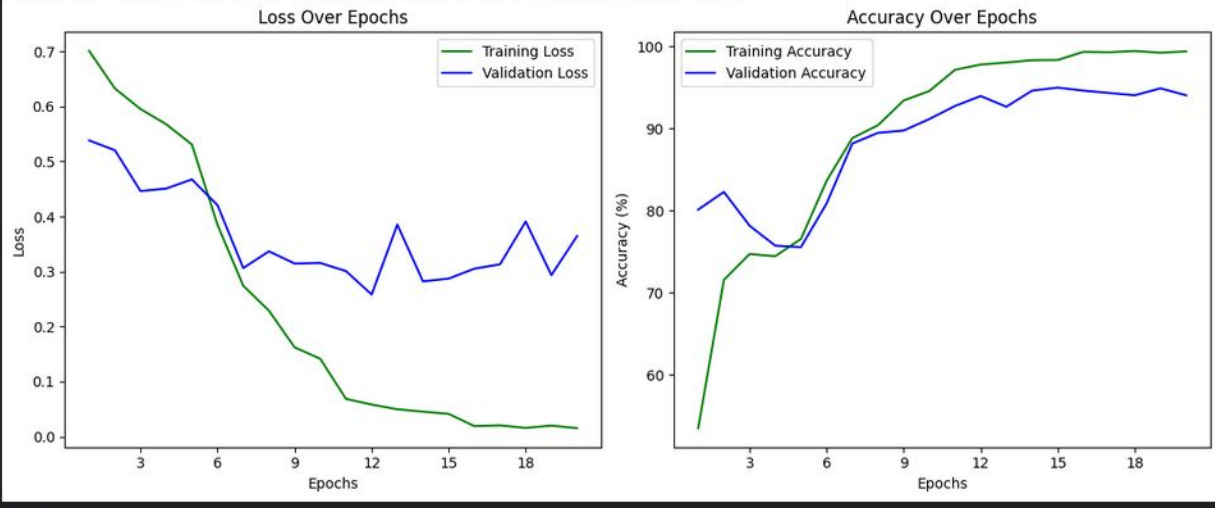
Training finished! Total Time: 224.652 min



# Experimental Results - SWIN

BS : 8  
num\_frames=10

Epoch 20/20, Train Loss: 0.0153, Train Acc: 99.36%, Val Loss: 0.3644, Val Acc: 94.02%



```
Confusion Matrix:
[[117  20]
 [ 13 387]]
Accuracy: 93.85%
```

```
Training finished! Total Time: 205.298 min
```

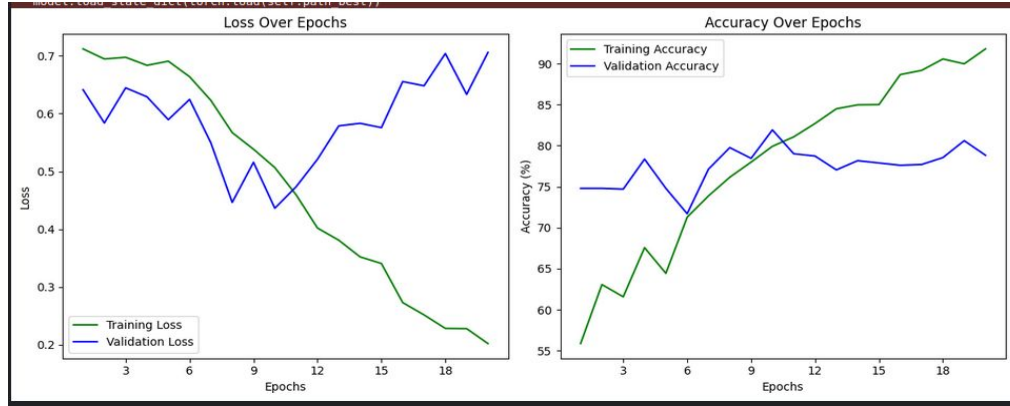
## Classification Report:

	precision	recall	f1-score	support
Real	0.9000	0.8540	0.8764	137
Fake	0.9509	0.9675	0.9591	400
accuracy			0.9385	537
macro avg	0.9254	0.9108	0.9178	537
weighted avg	0.9379	0.9385	0.9380	537

BS=8	F=10	93.85%
BS=4	F=20	75.19%

# Experimental Results - SWIN

BS : 4  
num\_frames=20



```
Confusion Matrix:
[[ 62  75]
 [ 58 341]]
Accuracy: 75.19%
```

```
Training curves saved to results/training_metrics.png
Training finished! Total Time: 403.409 min
```

```
Classification Report:
              precision    recall  f1-score   support

   Real         0.5167      0.4526      0.4825        137
   Fake         0.8197      0.8546      0.8368        399

   accuracy          0.7519
  macro avg         0.6682      0.6536      0.6597
 weighted avg         0.7423      0.7519      0.7462
```

Model	Parameters	Accuracy		Precision	Recall	F1 Score
EfficientNetB0	F=10 BS:8	94.79%	Real	0.90	0.92	0.91
			Fake	0.97	0.96	0.96
	F=20 BS:4	96.08%	Real	0.91	0.86	0.88
			Fake	0.95	0.97	0.96
	F=30 BS:4	96.76%	Real	0.94	0.92	0.93
			Fake	0.97	0.98	0.97
ResNext-50	F=10 BS:8	95.16%	Real	0.91	0.89	0.90
			Fake	0.96	0.97	0.96
	F=20 BS:4	95.52%	Real	0.88	0.94	0.91
			Fake	0.97	0.96	0.97
	F=30 BS:2	96.56%	Real	0.91	0.95	0.93
			Fake	0.98	0.97	0.98
SWIN	F=10 BS:8	93.85%	Real	0.90	0.85	0.87
			Fake	0.95	0.97	0.96
	F=20 BS:4	75.19%	Real	0.51	0.45	0.48
			Fake	0.81	0.85	0.84

Method	FF++ (LQ)
P3D [47]	67.05
R3D [67]	87.72
I3D [3]	93.18
M2TR <i>mean</i>	93.95
ST-M2TR	<b>95.31</b>

Paper: Multi-modal Multi-scale Transformers for Deepfake Detection (2022)

# Testing on other “unknown” datasets

Model	Parameters	Dataset		Precision	Recall	F1Score	Accuracy
ResNext-50	BS:2 F:10	UADFV	Real	0.825	0.97	0.89	0.88
			Fake	0.96	0.79	0.87	
		CelebDFv1	Real	0.53	0.90	0.67	0.70
			Fake	0.92	0.60	0.72	
	BS:2 F:20	UADFV	Real	0.84	0.82	0.83	0.83
			Fake	0.82	0.85	0.84	
		CelebDFv1	Real	0.51	0.90	0.65	0.67
			Fake	0.91	0.55	0.69	
	BS:2 F:30	UADFV	Real	0.96	0.93	0.95	0.95
			Fake	0.93	0.96	0.95	
		CelebDFv1	Real	0.47	0.97	0.63	0.61
			Fake	0.97	0.43	0.59	

Model	Parameters	Dataset		Precision	Recall	F1 Score	Accuracy
EfficientNetB0	BS:8 F:10	UADFV	Real	0.93	0.91	0.92	0.92
			Fake	0.91	0.94	0.92	
		CelebDFv1	Real	0.47	0.82	0.68	0.62
			Fake	0.85	0.52	0.65	
	BS:8 F:20	UADFV	Real	0.96	0.82	0.88	0.89
			Fake	0.84	0.97	0.94	
		CelebDFv1	Real	0.62	0.80	0.70	0.76
			Fake	0.88	0.74	0.80	
	BS:4 F:30	UADFV	Real	0.93	0.87	0.90	0.90
			Fake	0.88	0.93	0.90	
		CelebDFv1	Real	0.52	0.73	0.61	0.69
			Fake	0.83	0.66	0.74	

Model	Parameters	Dataset		Precision	Recall	F1 Score	Accuracy
SWIN	BS:2 F:10	UADFV	Real	1	0.85	0.92	0.92
			Fake	0.87	1	0.93	
		CelebDFV1	Real	0.47	0.70	0.56	0.63
			Fake	0.80	0.6	0.68	
	BS:2 F:20	UADFV	Real	0.96	0.85	0.90	0.91
			Fake	0.86	0.97	0.91	
		CelebDFv1	Real	0.51	0.80	0.62	0.67
			Fake	0.85	0.60	0.71	

UADFV : 95% with ResNext-LSTM\_F30  
CelebDFv1 : 76% with EffB0-LSTM\_F16

# Testing on other “unknown” datasets

Methods	Training data	UADFV [58]	Celeb-DF [31]
Two-stream [61]	Private data	85.1	55.7
Meso4 [6]	Private data	84.3	53.6
MesoInception4 [6]		82.1	49.6
HeadPose [58]	UADFV	89.0	54.8
FWA [30]	UADFV	97.4	53.8
VA-MLP [38]	Private data	70.2	48.8
VA-LogReg [38]		54.0	46.9
Multi-task [39]	FF	65.8	36.5
Xception-FF++ [42]	FF++	80.4	38.7

Paper: On the Detection of Digital Face Manipulation (2022)

Training Set	Testing Set	M2TR <sub>ncl</sub>	M2TR
FF++	Celeb-DF	65.6	68.2
	SR-DF	60.4	63.7

Paper: Multi-modal Multi-scale Transformers for Deepfake Detection (2022)

Generalisation Results for FF++ (HQ)

Training Set	Testing Set	Xception [49]	Multi-task [42]	Capsule [43]	DSW-FPA [35]	Two-Branch [39]	F3-Net [46]	MaDD [71]	DCViT [65]
FF++	FF++	99.7	76.3	96.6	93.0	98.7	98.1	99.3	98.3
	Celeb-DF	48.2	54.3	57.5	64.6	<b>73.4</b>	65.2	67.4	60.8

Table 1: Results on DFDC test

Model	AUC
ViT with distillation [18]	0.978
Selim EfficientNet B7 [37] <sup>†</sup>	0.972
Convolutional ViT [39]	0.843
Efficient ViT (our)	0.919
Conv. Cross ViT Wodajo CNN (our)	0.925
Conv. Cross ViT Eff.Net B0 - Avg (our)	0.947
Conv. Cross ViT Eff.Net B0 - Voting (our)	0.951

Model	Mean I
Convolutional ViT [39]	67%
Efficient ViT (our)	76%
Conv. Cross ViT Wodajo CNN (our)	76%

Model accuracy on FF++

Paper: Combining Efficient Net and Vision Transformers for Video Deepfake Detection (2022)


# Web GUI (Model Ensembling)





## Deepfake Detection System


Upload a video file, a zip file of a folder containing frames, or a single image. If multiple inputs are provided, the system will prioritize the folder of frames, then the image, and finally the video. The system extracts faces and classifies whether the input is real or fake using three different models with majority vote ensembling.

Input Video (optional)


  
Drop Video Here  
- or -  
Click to Upload




 

Upload Folder of Frames (zip file, optional)

  
Drop File Here  
- or -  
Click to Upload

Input Image (optional)

  
Drop Image Here  
- or -  
Click to Upload

Number of Frames to Process

Face Detection Confidence Threshold


0.9

1


Clear

Submit

Processed Faces (Annotated)



Prediction



Processing Details

Flag



# Web GUI (Model Ensembling)



## Deepfake Detection System

Upload a video file, a zip file of a folder containing frames, or a single image. If multiple inputs are provided, the system will prioritize the folder of frames, then the image, and finally the video. The system extracts faces and classifies whether the input is real or fake using three different models with majority vote ensembling.

The screenshot displays the user interface of the Deepfake Detection System. It includes an input section for video, folder of frames, or image, a processing area showing extracted faces, and a results section displaying the prediction (Fake) and processing details for three models: EfficientNet-B0+LSTM, ResNeXt+LSTM, and the final prediction (Fake).

**Input Video (optional)**

**Processed Faces (Annotated)**

**Prediction**

**Fake**

**EfficientNet-B0+LSTM:**  
Real: 0.3% | Fake: 99.7% | Prediction: Fake

**ResNeXt+LSTM:**  
Real: 0.0% | Fake: 100.0% | Prediction: Fake


# Web GUI (Model Ensembling)



### Deepfake Detection System

Upload a video file, a zip file of a folder containing frames, or a single image. If multiple inputs are provided, the system will prioritize the folder of frames, then the image, and finally the video. The system extracts faces and classifies whether the input is real or fake using three different models with majority vote ensembling.

Input Video (optional)




Upload Folder of Frames (zip file, optional)

Drop File Here  
- or -  
Click to Upload

Input Image (optional)

Drop Image Here  
- or -  
Click to Upload

Processed Faces (Annotated)



ii. Prediction

Real

Real ..... 100%

Fake ..... 0%


Processing Details

EfficientNet-B0+LSTM:  
Real: 99.9% | Fake: 0.1% | Prediction: Real

ResNeXt+LSTM:  
Real: 99.9% | Fake: 0.1% | Prediction: Real



# Demo





# Project Members

<u>Name</u>	<u>RollNo</u>	<u>Tasks</u>	<u>%</u>
Divyanshu	241110023	Dataset Gathering, Analysis & Preprocessing +UI	16
Khushwant	241110035	CNN+LSTM Deepfake Detection +UI & Generalisation Analysis	18
Krishanu	241110037	Hyperparameter Tuning & Testing of All Models,Noting Observations	16
Rishit	241110056	Image Deepfake Detection + UI	18
Rajan Kumar	241110087	SWIN Transformer Pipeline & Prediction Workflow	16
Senthil Ganesh	241110089	SWIN Transformer Pipeline & Prediction Workflow	16

Thank you!

